

VE527

Computer-Aided Design of Integrated Circuits

Maze Routing Implementation Issues

Outline

- Maze Routing Implementation Issues
 - Expansion Step
 - Data Structures
 - Depth-First Search
 - Global Routing

Implementation Concerns

- Algorithms
 - Question: if we can only process one cell at a time...
 - ...then, which cell is next to “label” in the search process
- Data structure (representation)
 - How do we store the routing grid?
 - What do we need in each cell?
 - How do we represent the state of the path search process?

First Issue: Expansion Step

- One big goal
 - Efficient storage: Big layout needs a big grid; put as little in each cell as possible.
- Big Idea: **Do not** store path costs in grid cells
 - If we have big path cost, then we need many bits per cell.
 - Notice: Only cells **most recently labeled** during search are used to **expand** the search.
 - Important terminology: these cells comprise the **search wavefront**.
 - It is the **wavefront** that is really most important....

Important Idea: Search Wavefront


- What is wavefront?
 - Set of cells already **reached** in the expansion process...
 - ...that **have neighbors** that we have **not** yet reached.
 - Wavefront is the **frontier** of the active search for new paths.
- Implication
 - **Don't** store the pathcost numbers **in the grid**.
 - Just store the wavefront cells themselves in a **separate data structure**.

	3	2	3		
3	2	1	2	3	
2	1	S	1	2	3
3	2	1	2	3	
	3	2	3		
		3			


How to Expand Wavefront?

Answer: Expand in the **pathcost** order

After expanding neighbors of
the cost “0” grid, i.e., S



After expanding neighbors of
the cost “1” grid




cost 1



cost 3

More Complex Wavefront

After expanding neighbors of the cost “2” grid



- **Expanded** in **pathcost order**, cheapest cells before more expensive cells.
- Algorithm: **Cheapest-cell-first search** (a variants of Dijkstra's algorithm)

Cheapest-Cell-First Search


- Find a **cheapest** cell C from the **wavefront**.
- Find the neighbors $\{N1, N2, \dots Nk\}$ of cell C that you have **not reached** yet.
- Compute the cost of expanding through C to reach the cells in $\{N1, N2, \dots Nk\}$.
- Add new cells $\{N1, N2, \dots Nk\}$ to the wavefront data structure, including their newly computed pathcosts, $\{\text{pathcost}(N1), \text{pathcost}(N2), \dots, \text{pathcost}(Nk)\}$.
 - Also, add to each of these new cells a pointer to cell C as the **predecessor**.
 - Mark the routing grid to remember that we have **reached** these cells $\{N1, N2, \dots Nk\}$.
- Remove cell C from the wavefront.
- Repeat with the **next cheapest cell on wavefront**...

Cell Predecessor


- **Predecessor** of cell A on the wavefront is a “tag” that tells **direction** from which cell A was reached, i.e., it’s how we found A.
- Why do we have to remember this?
 - Because we **do not mark pathcosts** in the grid any longer.
 - Then, to backtrace the path, we cannot simply follow the numbers in decreasing order.
 - We need a real “trail” that points, cell to cell, from target back to source.

Marking the Predecessor


- Assume 2 routing layers. 6 $\text{pred}(A)$ tags:
N, S, E, W, Up, Down
 - We only need 3 bits to store 6 tags.
 - **Note:** $\text{pred}(A)$ is the direction from which cell A was reached




$\text{prec}(A)=N$



$\text{prec}(A)=S$





$\text{prec}(A)=E$




$\text{prec}(A)=W$

Marking the Predecessor (cont.)




We Store Predecessors on Grid

Layer 1




Layer 2



- Result: backtrace is easy! Follow $\text{pred}(C)$ tags in the grid.

Example: Expanding & Reaching



- Expand cell B, reach cell A
 - B is with **cheapest** cost on wavefront.
 - A is B's **unreached** neighbor.
 - Compute cost to reach A from B:
 $12 + 3 = 15$
 - Add this new "cell object" to the wavefront:
 $\langle \text{layer} = L; \text{cell location} = (x_A, y_A); \text{pathcost} = 15 \rangle$
 - Mark cell A as $\text{Reached}(A) = \text{true}$
 - In future, we won't try to put it on the wavefront **again**.
 - In cell A in the grid, set $\text{pred}(A) = W$

Basic Maze Routing Algorithm

```
wavefront = { source cells };  
while (we have not reached target cell) {  
    if ( wavefront == empty ) quit; // no path to be found  
    C = get lowest cost cell on wavefront structure;  
    if ( C == target ) {  
        backtrace path in grid; // follow pred( ) pointers to source  
        cleanup & return; // we found a path  
    }  
    foreach ( unreached neighbor N of cell C ) {  
        mark cell N in grid as reached.  
        mark pred(N) with direction from cell N to C;  
        // compute cost to reach  
        pathcost(N) = pathcost(C) + cellcost(N) ;  
        add cell N to wavefront, indexed by pathcost(N);  
    }  
    delete cell C from wavefront;  
}
```

Observations


- Core algorithm is fairly simple. Works for grid with non uniform costs.
- Extends naturally to other features
 - **Multi-point nets**: just mark all the cells of intermediate paths as source, put them all in the wavefront to route to net point.
 - **Multiple routing layers**: just use parallel grids. When you expand to reach neighbors, check UP, DOWN to other layers. Use $\text{pred}(\text{Cell}) = \text{UP, DOWN}$ to do backtrace.

Outline

- Maze Routing Implementation Issues
 - Expansion Step
 - Data Structures
 - Depth-First Search
 - Global Routing


Two Key Data Structures: Overview

- Routing grid
 - Routing surface. Contains blockage. Has costs of each cell.
 - Mark these cells if they have already been reached.
 - Mark predecessor.
- Wavefront
 - Holds active cells **to be expanded**.
 - Cells store pathcost.
 - Indexed on pathcost; always expand cell with cheapest pathcost next.



Routing Grid

- Each grid cell C stores:
 - (x, y) coordinate of cell C
 - Layer of grid cell C
 - $\text{Cost}(C)$ (a small number)
 - **Note:** we can use $\text{cost} = -1$ to indicate blockage
 - $\text{Pred}(C)$ tag = N,S,E,W,Up,Down
 - $\text{Reached}(C)$, 1 bit Boolean true/false



Wavefront


- Each cell C in wavefront stores:
 - (x, y) coordinate of cell C
 - Layer of grid cell C
 - $\text{Pathcost}(C) = \text{length of the shortest path to cell } C$
- We need something clever here:
 - Need fast **extracting minimum** operations, to take a cell with the smallest pathcost from wavefront
 - Need fast insertion, to add a cell into the wavefront

Wavefront Structure: Min Heap

- Store cells of wavefront in a **min heap**.
 - Also called a **min priority queue**.
 - Classical data structure designed for fast insertion and extracting minimum operations.
 - These two operations have $O(\log N)$ time complexity for N objects.

Recall: Min Heap

- A min heap is a **complete binary tree** that
 - For **any** node v , the key of v is smaller than or equal to (\leq) the keys of any **descendants** of v .
- The item with the smallest cost is the **root**!
- Insert operation “bubbles up” the new item in heap to ensure the min heap property.
- Extracting minimum operation also ensures the min heap property.



Outline

- Maze Routing Implementation Issues
 - Expansion Step
 - Data Structures
 - Depth-First Search
 - Global Routing

Expansion Process, Revisited

- Problem:
 - Expand lots of cells to find one path to the target.
 - CPU time is proportional to number of cells you expand.
 - No attempt to search in direction of target first.
- Questions:
 - Can we actually bias expansion so we search **toward** the target?
 - Can we do this and still keep **guarantees** of reaching target with minimum cost path?

Motivation for Smart Search

Expanding **away** from the target seems to be a waste of time.


	3	2	3		
3	2	1	2	3	
2	1	S	1	2	3
3	2	1	2	3	
	3	2	3	T	
		3			

Searching **toward** the target in the shaded region, sometimes called the “source-target box”, seems a good idea to try first.

Smarter Search: Maze Search With Predictor

- Add **predictor function** to the cost, to direct the search **toward** the target.
- Plain maze router
 - Add a cell C to wavefront with cost that measures cost of **partial path**, $\text{source-to-cell}(C)$.
- Smarter maze router
 - Add cell C to wavefront with cost that **estimates entire source-to-target cost** of path **through the cell**.
 - Trick: estimate this as $\text{pathcost}(\text{source to cell } C) + \text{predictor}(\text{cell } C \text{ to target})$.
 - We call this method as smart depth first search (DFS) with a predictor in contrast to classical breadth first search (BFS).


Plain Maze Routing



Reached cell goes on wavefront with this cost:

$\text{pathcost}(\text{source to cell being expanded})$
+ cost of newly reached cell

Add A Depth-First Predictor



Reached cell goes on wavefront with this cost:

pathcost(source to cell being expanded)
+ cost of newly reached cell

+ estimate of pathcost(from newly reached cell to target)

A typical estimator is:

(Manhattan distance to target) ×
(Smallest cell cost)

Technical Results

- Depth first predictor
 - It is famous application of a classical idea: **A* search**.
 - Since predictor is a **lower bound** on extra pathcost you add to reach target...
 - ...by the property of A* search, you are **guaranteed** to get the minimum cost path.
- What does it do?
 - It **alters the order** in which we expand cells.
 - It prefers to expand cells that are **closer to the target first**.


Depth-First Expansion Example

Pathcost = 1;

Distance to target = 5

Esitmate of cost: $1+5=6$

Observation: Search prefers to stay **inside** the bounding box of the source-target rectangle before it expands other cells.



Pathcost = 1;

Distance to target = 3

Esitmate of cost: $1+3=4$


Depth-First Expansion Example

- It turns out for this example all the cells in this box cost the same!
 - **In this box:**
 $\text{pathcost}(\text{source} \rightarrow \text{cell}) + \text{estimated_pathcost}(\text{cell} \rightarrow \text{target}) = \text{constant}$
 - So – it expands **toward** the target!


		S	4	4	
		4	4	4	
		4	4	T	

DFS Expansion: Unintended Consequences

- This is “the” example of which “stay in the box” is **inefficient**.
 - The target is blocked inside the source-target box.
 - **Problem**: DFS explores the **whole shaded** rectangle before it tries anyplace else.
 - In this case, it may be faster to search outside the box!
- But, generally, this heuristic is good!



A Variation of the Heuristic



Reached cell goes on wavefront with this cost:

$\text{pathcost}(\text{source to cell being expanded})$
 $+ \text{cost of newly reached cell}$

$+ \boxed{K} \times \text{estimate of pathcost}(\text{from newly reached cell to target})$

K is usually $(1 + \text{a very small number})$


Idea: Insist that expand in a cheaper path toward target. Distorts optimum a bit, but it is faster.

Outline

- Maze Routing Implementation Issues
 - Expansion Step
 - Data Structures
 - Depth-First Search
 - Global Routing

Reality: ASIC Scale & Complexity


- Big chip, 1cm x 1cm
- 20-50 million nets
- Modern IC technology, $\sim 100\text{nm}$ pitch for wires (grid size)
- So, 100K x 100K routing grid!
- 10 routing layers
- 100 billion grid cells!
- Do we really do it like this?



No!


One Way: Preferred Routing Directions

- Every layer of metal wiring has a **preferred direction**.
 - E.g., metal 3, 5, 7, 9 prefer **vertical**; metal 4, 6, 8, 10 prefer **horizontal**.
 - As a result, all the wires on these layers look like straight lines.
 - If you need to bend, use a **via**. Makes it easier to embed lots of wires.
- This is one way to reduce complexity.




Divide & Conquer: Global Routing

- All the routing we have done so far is called **Detailed Routing**, which plans **exact, final** path of wire.
- How to deal with huge scale of chips? **Global Routing**
 - Global routing also imposes a grid on surface of the chip.
 - But now, this grid is much coarser, e.g., each grid cell is about 200 wire grids \times 200 wire grids in size.
- Big Idea: **Maze route** through these big, coarse regions.




Global vs Detailed Routing: Geometry



- Each cell in **Global Routing** grid is a box (called a **GBOX**) with size typically 100~200 wire grids on a side.
 - This example just has 6 for clarity.

What Global Routing Does?


- Global routing tells us we want this net to use this **rough** path; but not on which **exact** path inside this region.
 - Generate **regions of confinement** (i.e., coarse path) for a wire.
- Makes sure overall **wiring congestion** is reasonable.
 - Balance **supply** (how much available space) vs **demand** (how many paths want to go here).



What Detailed Routing Does?

- Detailed routing embeds **exact** paths in these regions.
 - Simplest model: **a finer grid**. Require wires to use tracks on this grid.


This reduces computation!



Global Routing as Maze Routing

- What's different for a global router? **Cell Cost Function**
 - **Detailed router**: grid cells are blocked (cost=infinity) or they have a cost proportional to how much we want to avoid them.
 - **Global router**: grid cells have **dynamic** cost, proportional to how many **previously** routed wires used this cell, how many more can fit **later**.

Suppose global GBOX is 10 tracks for wires.



Maze Routing: Summary

- Has been around a long time.
 - Very **flexible** cost-based search; can be recast to attack many problems.
 - **Dominates** most modern routers...
 - Some other approaches, e.g., sophisticated "grid-less" representation of "space".
- Lots of ancillary problems (and solutions)
 - Hierarchy/abstraction (e.g., global routing).
 - Iterative improvement strategies.